

# Week 5: Assignment 4: Word relationship analysis

Scout Leonard

2022-05-02

## Load Libraries

```
packages=c("tidyr",
            "pdftools",
            "lubridate",
            "tidyverse",
            "tidytext",
            "readr",
            "quanteda",
            "readtext",
            "quanteda.textstats",
            "quanteda.textplots",
            "ggplot2",
            "forcats",
            "stringr",
            "quanteda.textplots",
            "widyrr",
            "igraph",
            "ggraph",
            "here")

for (i in packages) {
  if (require(i,character.only=TRUE)==FALSE) {
    install.packages(i,repos='http://cran.us.r-project.org')
  }
  else {
    require(i,character.only=TRUE)
  }
}
```

## Read in data

```
#filepath to data
files <- list.files(path = here("data/week5_data/"),
                    pattern = "*.pdf$",
                    full.names = TRUE)

#renders all textboxes on a text canvas and returns a character vector of equal length to the number of
ej_reports <- lapply(files, pdf_text)
```

```

#read texts and (if any) associated document-level meta-data from one or more source files - makes a df
ej_pdf <- readtext(file = here("data/week5_data/*.pdf"),
                  docvarsfrom = "filenames",
                  docvarnames = c("type", "subj", "year"),
                  sep = "_")

#creating an initial corpus containing our data
epa_corp <- corpus(x = ej_pdf, text_field = "text" )

#return details of the corpus
summary(epa_corp) %>%
  knitr::kable()

```

Text	Types	Tokens	Sentences	type	subj	year
EPA_EJ_2015.pdf	2136	8944	263	EPA	EJ	2015
EPA_EJ_2016.pdf	1599	7965	176	EPA	EJ	2016
EPA_EJ_2017.pdf	2774	16658	447	EPA	EJ	2017
EPA_EJ_2018.pdf	3973	30564	653	EPA	EJ	2018
EPA_EJ_2019.pdf	3773	22648	672	EPA	EJ	2019
EPA_EJ_2020.pdf	4493	30523	987	EPA	EJ	2020

```

#I'm adding some additional, context-specific stop words to stop word lexicon
more_stops <-c("2015", "2016", "2017", "2018",
              "2019", "2020", "www.epa.gov", "https")

#add the additional stopwords to the stop word lexicon
add_stops <- tibble(word = c(stop_words$word, more_stops))

stop_vec <- as_vector(add_stops)

#convert to tidy format and apply my stop words
raw_text <- tidy(epa_corp)

#Distribution of most frequent words across documents
raw_words <- raw_text %>%
  mutate(year = as.factor(year)) %>%
  unnest_tokens(word, text) %>%
  anti_join(add_stops, by = 'word') %>%
  count(year, word, sort = TRUE)

#number of total words by document
total_words <- raw_words %>%
  group_by(year) %>%
  summarize(total = sum(n))

report_words <- left_join(raw_words, total_words)

## Joining, by = "year"
par_tokens <- unnest_tokens(raw_text,
                          output = paragraphs,
                          input = text,
                          token = "paragraphs")

```

```
par_tokens <- par_tokens %>%  
  mutate(par_id = 1:n())  
  
par_words <- unnest_tokens(par_tokens,  
                           output = word,  
                           input = paragraphs,  
                           token = "words")
```

## Part 1

What are the most frequent trigrams in the dataset? How does this compare to the most frequent bigrams? Which n-gram seems more informative here, and why?

```
tokens <- tokens(epa_corp, remove_punct = TRUE) #list of character vectors - takes each document and sp

toks1<- tokens_select(tokens, min_nchar = 3)

toks1 <- tokens_tolower(toks1)

toks1 <- tokens_remove(toks1, pattern = (stop_vec))

dfm <- dfm(toks1) #create document feature matrix - rows are number of occurrences of each word within e

#first the basic frequency stat
tstat_freq <- textstat_frequency(dfm, n = 5, groups = year)

head(tstat_freq, 10) %>%
  knitr::kable()
```

feature	frequency	rank	docfreq	group
environmental	127	1	1	2015
communities	99	2	1	2015
epa	92	3	1	2015
justice	84	4	1	2015
community	47	5	1	2015
environmental	109	1	1	2016
communities	85	2	1	2016
justice	71	3	1	2016
epa	48	4	1	2016
federal	31	5	1	2016

```
toks2 <- tokens_ngrams(toks1, n = 3)

dfm2 <- dfm(toks2)

dfm2 <- dfm_remove(dfm2, pattern = c(stop_vec))
#gives more coherent terms - power of chunking at a different token level

freq_words2 <- textstat_frequency(dfm2, n = 20)

freq_words2$token <- rep("trigram", 20)
#tokens1 <- tokens_select(tokens1,pattern = stopwords("en"), selection = "remove")

head(freq_words2, 5) %>%
  knitr::kable()
```

feature	frequency	rank	docfreq	group	token
justice_fy2017_progress	51	1	1	all	trigram
fy2017_progress_report	51	1	1	all	trigram
environmental_public_health	50	3	6	all	trigram
environmental_justice_fy2017	50	3	1	all	trigram

feature	frequency	rank	docfreq	group	token
national_environmental_justice	37	5	6	all	trigram

The most frequent trigrams in the dataset are shown in the table above, with justice\_fy2017\_progress as the most frequently occurring trigram.

## Part 2

Choose a new focal term to replace “justice” and recreate the correlation table and network (see `corr_paragraphs` and `corr_network` chunks). Explore some of the plotting parameters in the `cor_network` chunk to see if you can improve the clarity or amount of information your plot conveys. Make sure to use a different color for the ties!

```
#correlation between co-occurring words
word_cors <- par_words %>%
  add_count(par_id) %>%
  filter(n >= 50) %>%
  select(-n) %>%
  pairwise_cor(word, par_id, sort = TRUE)

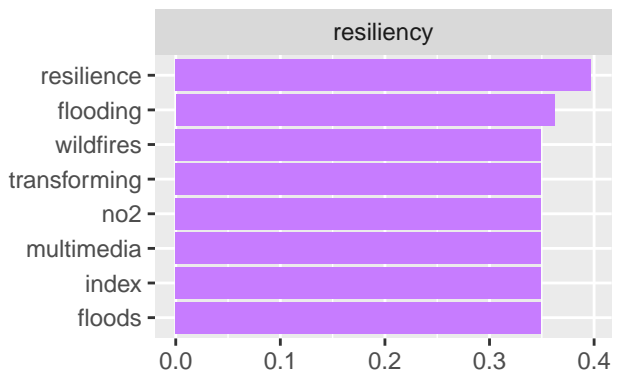

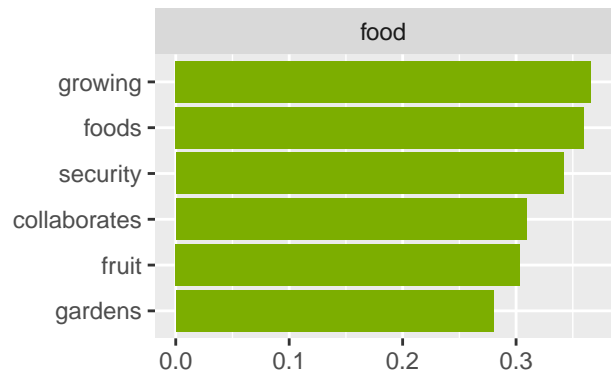
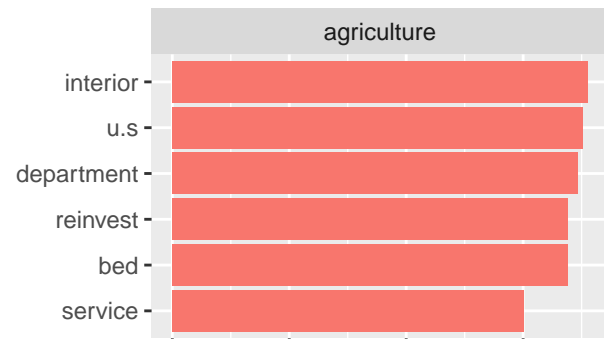
#now we can select words cooccurring with the word justice and get correlation coefficients
food_cors <- word_cors %>%
  filter(item1 == "food")

word_cors %>%
  filter(item1 %in% c("food", "agriculture", "health", "resiliency")) %>%
  group_by(item1) %>%
  top_n(6) %>%
  ungroup() %>%
  mutate(item1 = as.factor(item1),
         name = reorder_within(item2, correlation, item1)) %>%
  ggplot(aes(y = name, x = correlation, fill = item1)) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~item1, ncol = 2, scales = "free")+
  scale_y_reordered() +
  labs(y = NULL,
       x = NULL,
       title = "Correlations with key words",
       subtitle = "EPA EJ Reports")
```

```
## Selecting by correlation
```

## Correlations with key words

### EPA EJ Reports



```
#let's zoom in on just one of our key terms
food_cors <- word_cors %>%
  filter(item1 == "food") %>%
  mutate(n = 1:n())
```

```
food_cors %>%
  filter(n <= 50) %>%
  graph_from_data_frame() %>%
  ggraph(layout = "fr") +
  geom_edge_link(aes(edge_alpha = correlation,
                     edge_width = correlation),
                edge_colour = "darkmagenta") +
  geom_node_point(size = 3) +
  geom_node_text(aes(label = name),
                repel = TRUE,
                point.padding = unit(0.2,
                                   "lines")) +
  theme_void()
```





## Part 3

Write a function that allows you to conduct a keyness analysis to compare two individual EPA reports (hint: that means target and reference need to both be individual reports). Run the function on 3 pairs of reports, generating 3 keyness plots.

```
#test function
```

```
dfm_subset <- corpus_subset(epa_corp, grepl("2018|2019", docnames(epa_corp)))
```

```
#write function
```

```
keyness_comparison <- function(text1_year, text2_year) {
```

```
  #subset the corpus
```

```
  corpus_subset <- corpus_subset(epa_corp, grepl(paste0(text1_year, "|", text2_year), docnames(epa_corp)))
```

```
  #tokenize corpus
```

```
  tokens <- tokens(corpus_subset, remove_punct = TRUE) #list of character vectors - takes each document
```

```
  toks <- tokens_select(tokens, min_nchar = 3)
```

```
  toks <- tokens_tolower(toks)
```

```
  toks <- tokens_remove(toks, pattern = (stop_vec))
```

```
  dfm <- dfm(toks) #create document feature matrix - rows are number of occurrences of each word within
```

```
  keyness <- textstat_keyness(dfm, target = 2)
```


```
  textplot_keyness(keyness)
```

```
}
```

## Keyness plot 1

Test running the function on 2015 and 2019


```
keyness_comparison(text1_year = 2015, text2_year = 2019)
```



## Keyness plot 2

Test running the function on 2019 and 2020


```
keyness_comparison(text1_year = 2019, text2_year = 2020)
```



### Keyness plot 3

Test running the function on 2016 and 2018:

```
keyness_comparison(text1_year = 2016, text2_year = 2018)
```



## Part 4

Select a word or multi-word term of interest and identify words related to it using windowing and keyness comparison. To do this you will create two objects: one containing all words occurring within a 10-word window of your term of interest, and the second object containing all other words. Then run a keyness comparison on these objects. Which one is the target, and which the reference? Hint

### **food systems**

Create an object containing all words occurring within a 10 word window of **food systems**.

```
term = "food"

toks_inside <- tokens_keep(toks1,
                           pattern = term,
                           window = 10)

toks_outside <- tokens_remove(toks1,
                              pattern = term)
```

Create an object containing all other words:

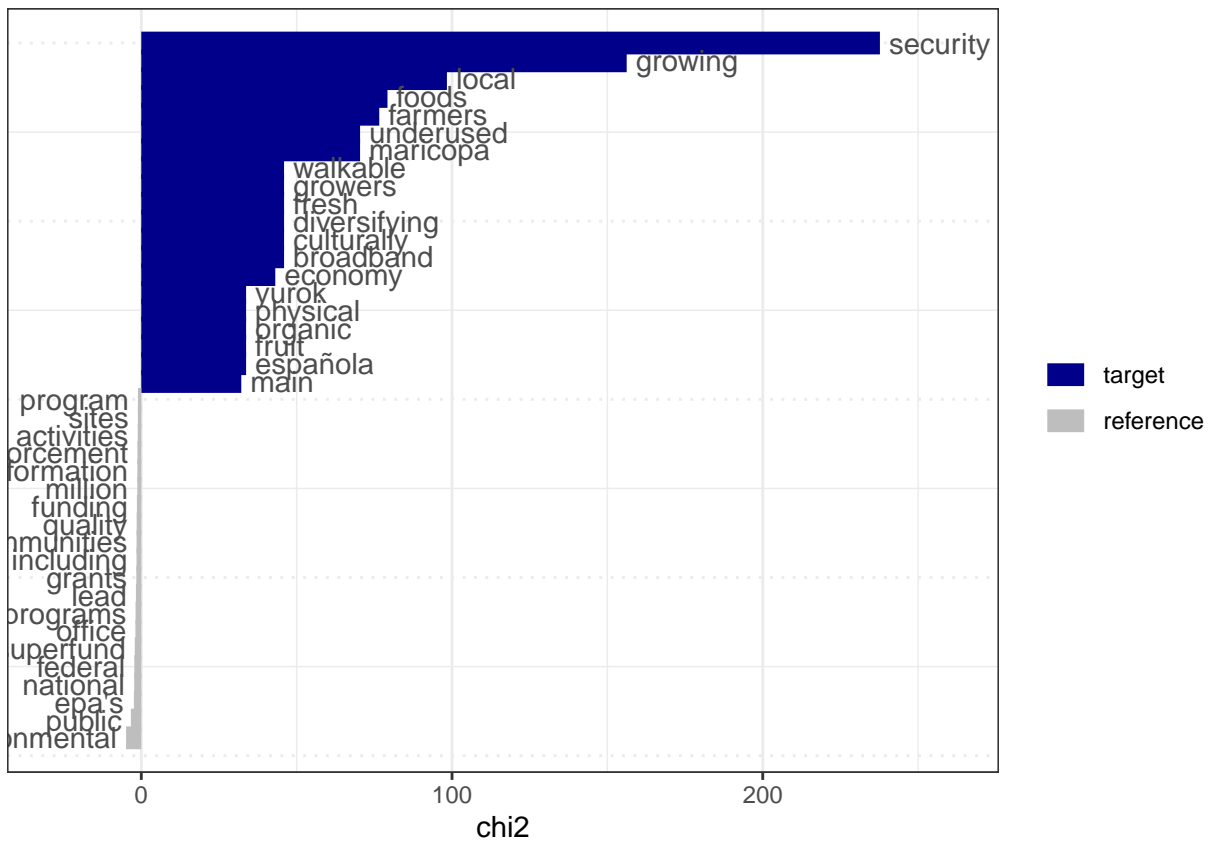
```
toks_outside <- tokens_remove(toks1,
                              pattern = term,
                              window = 10)
```

Run a keyness comparison of the objects:

```
dfmat_inside <- dfm(toks_inside)
dfmat_outside <- dfm(toks_outside)

tstat_key_inside <- textstat_keyness(rbind(dfmat_inside, dfmat_outside),
                                     target = seq_len(ndoc(dfmat_inside)))

textplot_keyness(tstat_key_inside)
```



`toks_inside` is the target, and `toks_outside` is the reference.